

SSW 590

Version (CL , SHA , )

by

Gavin Lam, Spurthi Setty, Annanya Jain, Luo Xu

glam1@Stevens.edu, lxu41@stevens.edu

October 16, 2025

© Gavin Lam, Spurthi Setty, Annanya Jain, Luo Xu  
glam1@Stevens.edu, lxu41@stevens.edu  
ALL RIGHTS RESERVED

SSW 590  
Version (CL , SHA , )

Gavin Lam, Spurthi Setty, Annanya Jain, Luo Xu  
glam1@Stevens.edu, lxu41@stevens.edu

This document provides the requirements and design details of the PROJECT. The following table (Table 1) should be updated by authors whenever major changes are made to the architecture design or new components are added. Add updates to the top of the table. Most recent changes to the document should be seen first and the oldest last.

Table 1: Document Update History

Date	Updates
10/11/2025	Luo: <ul style="list-style-type: none"><li>• Updated Overleaf Chapter to include all packages Overleaf requires (9)</li></ul>
10/05/2025	Annanya: <ul style="list-style-type: none"><li>• Updated and documented the Overleaf Chapter(9)</li></ul>
10/05/2025	Gavin: <ul style="list-style-type: none"><li>• Updated the Hosts and Password tables (2)(1)</li></ul>
10/03/2025	Luo: <ul style="list-style-type: none"><li>• Updated and documented the Bugzilla Chapter (8)</li></ul>
10/02/2025	Luo: <ul style="list-style-type: none"><li>• Created the Bugzilla (8)</li></ul>
10/01/2025	Annanya: <ul style="list-style-type: none"><li>• Added to chapter LaTeX Docker regarding the steps taken to create Docker container which compiles a simple latex file. (7)</li></ul>
09/29/2025	Gavin: <ul style="list-style-type: none"><li>• Added to chapter AWS the steps we took to get our website deployed. (6)</li></ul>
09/24/2025	Luo, Ananaya, Gavin: <ul style="list-style-type: none"><li>• Improved and added more to the Project Proposal in terms of tools, devop tools, and description. (5)</li><li>• Created the AWS Deployment file. (6)</li></ul>
09/21/2025	Luo: <ul style="list-style-type: none"><li>• Created the Project Proposal (5)</li></ul> Spurthi: <ul style="list-style-type: none"><li>• The Betrayal (she switched teams)</li></ul>

Table 1: Document Update History

Date	Updates
09/15/2025	Gavin, Spurthi, Annanya, Luo: <ul style="list-style-type: none"> <li>Completed remaining Linux Problem Sets in Linux Commands chapter (4)</li> </ul>
09/14/2025	Gavin, Spurthi, Annanya, Luo: <ul style="list-style-type: none"> <li>Added new chapter, Linux Commands. (4)</li> <li>Completed terminal session and solved A-D in the problem set (4)</li> </ul>
09/9/2025	Gavin, Spurthi, Annanya, Luo: <ul style="list-style-type: none"> <li>Updated Passwords chapter (Chapter 1) with long table with user/password/server rules.</li> <li>Created Hosts chapter (Chapter 2) and Kanban Setup (Chapter 3).</li> </ul>

# Table of Contents

<b>1</b>	<b>Passwords</b>	
	– <i>Gavin Lam, Spurthi Setty, Annanya Jain, Luo Xu</i>	<b>1</b>
<b>2</b>	<b>Hosts</b>	
	– <i>Gavin Lam, Spurthi Setty, Annanya Jain, Luo Xu</i>	<b>2</b>
<b>3</b>	<b>Kanban Setup</b>	
	– <i>Gavin Lam, Spurthi Setty, Annanya Jain, Luo Xu</i>	<b>3</b>
<b>4</b>	<b>Linux Commands</b>	
	– <i>Spurthi Setty, Gavin Lam, Annanya Jain, Luo Xu</i>	<b>7</b>
4.1	Terminal Session . . . . .	7
4.2	Problem-Set Commands and Outputs . . . . .	7
4.2.1	Navigation & File Ops . . . . .	7
4.2.2	Viewing & Searching . . . . .	8
4.2.3	Text Processing . . . . .	9
4.2.4	Permissions & Ownership . . . . .	10
4.2.5	Links & Find . . . . .	10
4.2.6	Processes & Job Control . . . . .	10
4.2.7	Archiving & Compression . . . . .	11
4.2.8	Networking & System Info . . . . .	12
4.2.9	Package & Services (Debian/Ubuntu) . . . . .	12
4.2.10	Bash & Scripting . . . . .	13
<b>5</b>	<b>Project Proposal</b>	
	– <i>Luo Xu, Annanya Jain, Gavin Lam</i>	<b>15</b>
<b>6</b>	<b>AWS</b>	
	– <i>Gavin Lam, Luo Xu, Annanya Jain</i>	<b>16</b>
<b>7</b>	<b>LaTeX Docker</b>	
	– <i>Annanya Jain, Gavin Lam, Luo Xu</i>	<b>24</b>
7.1	Steps taken for creating the docker container to compile a LaTeX File . . . . .	24

7.1.1	main.tex . . . . .	24
7.1.2	Dockerfile . . . . .	25
7.1.3	Building the Docker image . . . . .	25
<b>8</b>	<b>Bugzilla</b>	
	– <i>Luo Xu, Gavin Lam, Annanya Jain</i>	<b>28</b>
8.1	DigitalOcean Setup . . . . .	28
8.2	Bugzilla . . . . .	28
<b>9</b>	<b>Overleaf</b>	
	– <i>Annanya Jain, Luo Xu, Gavin Lam</i>	<b>33</b>
9.1	Overview . . . . .	33
9.2	Initial Access Issue . . . . .	33
9.3	Overleaf Setup Process . . . . .	34
9.3.1	Step 1: Create Project Directory . . . . .	34
9.3.2	Step 2: Create Environment File . . . . .	34
9.3.3	Step 3: Write the Docker Compose File . . . . .	34
9.3.4	Step 4: Launch the Containers . . . . .	35
9.4	Debugging the Overleaf Restart Loop . . . . .	35
9.4.1	Step 1: Check Logs . . . . .	35
9.4.2	Step 2: Error Observed . . . . .	35
9.5	Fixing MongoDB Configuration . . . . .	35
9.5.1	Step 1: Enter Mongo Shell . . . . .	35
9.5.2	Step 2: Initialize Replica Set . . . . .	37
9.5.3	Step 3: Update Feature Compatibility Version . . . . .	37
9.5.4	Step 4: Restart Overleaf . . . . .	37
9.6	Making Overleaf Publicly Accessible . . . . .	37
9.7	Access and Verification . . . . .	37
9.8	Conclusion . . . . .	38
<b>A</b>	<b>Appendix</b>	
	– <i>Gavin Lam, Spurthi Setty, Annanya Jain, Luo Xu</i>	<b>41</b>
	<b>Bibliography</b>	<b>43</b>

# List of Tables

1	Document Update History	iii
1	Document Update History	iv
1.1	Password Rules	1

# List of Figures

8.1	bugzilla website on <a href="http://174.138.68.199/">http://174.138.68.199/</a> . . . . .	31
9.1	overleaf website on <a href="http://174.138.68.199:8090">http://174.138.68.199:8090</a> . . . . .	39



# Chapter 1

## Passwords

– Gavin Lam, Spurthi Setty, Annanya Jain, Luo Xu

Table 1.1: Password Rules

User	Password Rule / Hint
OVERLEAF_ADMIN_EMAIL	Password should include a mix of upper and lower case letters, numbers, and a special character. Hint: key + specialcharacters
MONGO_INITDB_ROOT_USERNAME	Strong password required — must not contain common words or personal information. Hint: Secure passphrase based on your project title.
MONGO_INITDB_ROOT_PASSWORD	At least 10 characters, must include at least one uppercase, lowercase, number, and symbol. Hint: Prefix "Mongo" + symbol + random digits
bugzillauser (BugzillaDB)	Must include a combination of regular characters, numbers, and special characters. Hint: key + numbers + specialcharacters
bugzilla	Must include a combination of regular characters, numbers, and special characters. Hint: key + numbers + specialcharacters
devuser	At least 8 characters, include uppercase, lowercase, a number, and a symbol. Hint: First pet.
admin	Must change passwords every 90 days. Hint: Favorite City.
tester	Must include word banana Hint: Popular desert item.

# Chapter 2

## Hosts

– Gavin Lam, Spurthi Setty, Annanya Jain, Luo Xu

Host Name	Description
bugzilla	Bugzilla web application container running on port 80:80 to manage and track software bugs.
overleaf	Overleaf collaborative LaTeX editor container, accessible via port 8090. Configured to use MongoDB and Redis.
mongo	MongoDB service used by Overleaf for document and project data storage. Runs internally on port 27017.
redis	Redis in-memory cache used by Overleaf for session management and performance optimization.
digitalocean droplet	Shared host (Ubuntu) where both Bugzilla and Overleaf Docker environments are deployed. Accessible via SSH key linked to GitHub for secure root access.
bugzillaDB	Bugzilla database
bugzilla	Bugzilla server to catch bugs
dev-server	Primary development server.
test-server	Server for automated testing. scripts.
prod-server	Final server for working product.

# Chapter 3


## Kanban Setup

– Gavin Lam, Spurthi Setty, Annanya Jain, Luo Xu

For my DevOps project, I chose to use Atlassian JIRA to set up my Kanban board because I have some familiarity with it having used it once before in another class.

First, I created a new project in Jira and selected the **Kanban template**. I decided to use the team-managed project option and gave it the name SSW 590 after the class. JIRA automatically provided the columns **To Do**, **In Progress**, and **Done**.

### Step 1: Set up Jira Team




I have chosen the name 'stevens-ssw590-fall2025' as the site name.

## Step 2: Set up a Kanban Project



## Step 3: Kanban



png/kanban\_example.png

# Chapter 4

## Linux Commands

– Spurthi Setty, Gavin Lam, Annanya Jain, Luo Xu

### 4.1 Terminal Session

The following commands were run from ~/Documents/Devops/Assignment2 and create several test files and directories.

```
ubuntu@Ubuntu:~/Documents/Devops/Assignment2$ mkdir -p ~/lx-test && cd ~/lx-test
ubuntu@Ubuntu:~/lx-test$ printf "alpha\nbeta\nGamma\ngamma\nbeta\n" > words.txt
ubuntu@Ubuntu:~/lx-test$ printf
  ↪ "id,name,dept\n1,Ada,EE\n2,Linus,CS\n3,Grace,EE\n4,Dennis,CS\n" > people.csv
ubuntu@Ubuntu:~/lx-test$ printf "INFO boot ok\nWARN disk low\nERROR fan
  ↪ fail\nINFO shutdown\n" > sys.log
ubuntu@Ubuntu:~/lx-test$ dd if=/dev/zero of=blob.bin bs=1K count=48 status=none
ubuntu@Ubuntu:~/lx-test$ mkdir -p src/lib tmp archive
ubuntu@Ubuntu:~/lx-test$ printf "one two three four\n" > src/file1.txt
ubuntu@Ubuntu:~/lx-test$ printf "two three four five\n" > src/file2.txt
ubuntu@Ubuntu:~/lx-test$ ln -s src/file1.txt link-to-file1
ubuntu@Ubuntu:~/lx-test$ touch -t 202401020304 old.txt
```

### 4.2 Problem-Set Commands and Outputs

#### 4.2.1 Navigation & File Ops

1. Present working directory

```
ubuntu@Ubuntu:~/lx-test$ pwd
/home/ubuntu/lx-test
```

2. List all entries, including dotfiles

```
ubuntu@Ubuntu:~/lx-test$ ls -A1
archive
```

```
blob.bin
link-to-file1
old.txt
people.csv
src
sys.log
tmp
words.txt
```

3. **Copy src/file1.txt to tmp/ only if tmp exists (verbose)**

```
ubuntu@Ubuntu:~/lx-test$ test -d tmp && cp -v src/file1.txt tmp/
'src/file1.txt' -> 'tmp/file1.txt'
```

4. **Move old.txt into archive/ and keep timestamp**

```
ubuntu@Ubuntu:~/lx-test$ mv -v old.txt archive/
renamed 'old.txt' -> 'archive/old.txt'
```

5. **Create an empty notes.md only if it does not exist**

```
ubuntu@Ubuntu:~/lx-test$ [ -e notes.md ] || : > notes.md
```

6. **Show disk usage (human-readable) for src directory**

```
ubuntu@Ubuntu:~/lx-test$ du -sh src
16K    src
```

## 4.2.2 Viewing & Searching

7. **Print line numbers while displaying sys.log**

```
ubuntu@Ubuntu:~/lx-test$ nl sys.log
 1  INFO boot ok
 2  WARN disk low
 3  ERROR fan fail
 4  INFO shutdown
```

8. **Show only the lines in sys.log that contain ERROR (case-sensitive)**

```
ubuntu@Ubuntu:~/lx-test$ grep 'ERROR' sys.log
ERROR fan fail
```

9. **Count how many distinct words appear in words.txt (case-insensitive)**

```
ubuntu@Ubuntu:~/lx-test$ tr '[:upper:]' '[:lower:]' < words.txt | tr -s
↪ '[:space:]' '\n' | sort -u | wc -l
3
```



**10. From words.txt, show lines that start with g or G**

```
ubuntu@Ubuntu:~/lx-test$ grep -E '^[gG]' words.txt
Gamma
gamma
```

**11. Display the first 2 lines of people.csv without using an editor**

```
ubuntu@Ubuntu:~/lx-test$ head -n 2 people.csv
id,name,dept
1,Ada,EE
```

**12. Show the last 3 lines of sys.log and keep following if the file grows**

```
ubuntu@Ubuntu:~/lx-test$ tail -n 3 -f sys.log
WARN disk low
ERROR fan fail
INFO shutdown
```

### 4.2.3 Text Processing

**13. From people.csv, print only the name column (2nd), excluding the header.**

```
ubuntu@Ubuntu:~/lx-test$ tail -n +2 people.csv | cut -d',' -f2
Ada
Linus
Grace
Dennis
```

**14. Sort words.txt case-insensitively and remove duplicates**

```
ubuntu@Ubuntu:~/lx-test$ sort -f words.txt | uniq -i
alpha
beta
Gamma
```

**15. Replace every three with 3 in all files under src/ in-place, creating .bak backups.**

```
ubuntu@Ubuntu:~/lx-test$ find src -type f -exec sed -i.bak 's/three/3/g' {}
↪ +
```

**16. Print the number of lines, words, and bytes for every \*.txt file in src/.**

```
ubuntu@Ubuntu:~/lx-test$ wc src/*.txt
 1  4 15 src/file1.txt
 1  4 16 src/file2.txt
 2  8 31 total
```

## 4.2.4 Permissions & Ownership

17. Make tmp/ readable, writable, and searchable only by the owner.

```
ubuntu@Ubuntu:~/lx-test$ chmod 700 tmp/
```

18. Give group execute permission to src/lib recursively without touching others/owner bits.

```
ubuntu@Ubuntu:~/lx-test$ chmod -R g+x src/lib
```

19. Show the numeric (octal) permissions of src/file2.txt

```
ubuntu@Ubuntu:~/lx-test$ stat -c "%a" src/file2.txt  
664
```

20. Make notes.md append-only for the owner via file attributes (if supported).

```
ubuntu@Ubuntu:~/lx-test$ sudo chattr +a notes.md
```

## 4.2.5 Links & Find

21. Verify whether link-to-file1 is a symlink and show its target path.

```
luo@ubuntuluo:~/lx-test$ ls -l link-to-file1  
lrwxrwxrwx 1 luo luo 13 Sep 16 18:56 link-to-file1 -> src/file1.txt
```

22. Find all regular files under the current tree larger than 40 KiB.

```
luo@ubuntuluo:~/lx-test$ find . -type f -size +40k
```

23. Find files modified in the last 10 minutes under tmp/ and print their sizes.

```
luo@ubuntuluo:~/lx-test$ find tmp/ -type f -mmin -10 -exec ls -lh {} +
```

## 4.2.6 Processes & Job Control

24. Show your processes in a tree view.

```
luo@ubuntuluo:~/lx-test$ pstree -p
```

25. Start sleep 120 in the background and show its PID.

```
luo@ubuntuluo:~/lx-test$ sleep 120 &
echo $!
[1] 4474
4474
```

26. Send a TERM signal to all sleep processes owned by you (don't use kill -9).

```
luo@ubuntuluo:~/lx-test$ pkill -TERM -u "$USER" sleep

[1]  + Terminated  sleep 120
```

27. Show the top 5 processes by memory usage (one-shot, not interactive).

```
luo@ubuntuluo:~/lx-test$ ps -eo pid,ppid,user,%mem,%cpu,comm --sort=-%mem |
→ head -n 5
PID      PPID USER      %MEM %CPU COMMAND
  1925    1711 luo         9.8  4.0 gnome-shell
  2451    1925 luo         2.4  0.0 mutter-x11-fram
  2328    1711 luo         2.0  0.0 gsd-xsettings
  2245    1925 luo         1.7  0.0 Xwayland
```

## 4.2.7 Archiving & Compression

28. Create a gzipped tar archive src.tgz from src/ with relative paths.

```
luo@ubuntuluo:~/lx-test$ tar -czf src.tgz -C src .
```

29. List the contents of src.tgz without extracting.

```
luo@ubuntuluo:~/lx-test$ tar -tzf src.tgz
./
./file2.txt
./lib/
./file1.txt.bak
./file2.txt.bak
./file1.txt
```

30. Extract only file2.txt from src.tgz into tmp/.

```
luo@ubuntuluo:~/lx-test$ tar -xvzf src.tgz -C tmp ./file2.txt
./file2.txt
```

## 4.2.8 Networking & System Info

31. Show all listening TCP sockets with associated PIDs (no root assumptions).

```
annanyajain@ubuntu:~/lx-test$ ss -tlnp
State  Recv-Q  Send-Q    Local Address:Port    Peer Address:Port    Process
```

32. Print your default route (gateway) in a concise form.

```
annanyajain@ubuntu:~/lx-test$ ip route show default
default via 198.19.249.1 dev eth0 proto dhcp src 198.19.249.228 metric 100
```

33. Display kernel name, release, and machine architecture.

```
annanyajain@ubuntu:~/lx-test$ uname -srm
Linux 6.12.10-orbstack-00297-gf8f6e015b993 aarch64
```

34. Show the last 5 successful logins (or last sessions) on the system.

```
annanyajain@ubuntu:~/lx-test$ last -n 5
reboot    system boot  6.12.10-orbstack Tue Sep 16 01:01    still running
reboot    system boot  6.12.10-orbstack Tue Jan 21 13:27 - 13:40    (00:12)
reboot    system boot  6.12.10-orbstack Tue Jan 21 10:51 - 10:56    (00:04)
reboot    system boot  6.12.9-orbstack- Mon Jan 20 19:48 - 19:48    (00:00)
reboot    system boot  6.12.9-orbstack- Mon Jan 20 17:30 - 19:48    (02:18)

wtmp begins Mon Jan 20 15:07:29 2025
```

## 4.2.9 Package & Services (Debian/Ubuntu)

35. Show the installed version of package coreutils.

```
annanyajain@ubuntu:~/lx-test$ dpkg -s coreutils | grep '^Version:'
Version: 8.32-4.1ubuntu1.2
```

36. Search available packages whose names contain ripgrep.

```
annanyajain@ubuntu:~/lx-test$ apt search ripgrep
Sorting... Done
Full Text Search... Done
elpa-dumb-jump/jammy 0.5.3-1 all
  jump to definition for multiple languages without configuration

ripgrep/jammy-updates,jammy-security 13.0.0-2ubuntu0.1 arm64
  Recursively searches directories for a regex pattern
```

```
ugrep/jammy 3.7.2+dfsg-1 arm64  
faster grep with an interactive query UI
```

37. Check whether service cron is active and print its status line only.

```
annanyajain@ubuntu:~/lx-test$ systemctl status cron | grep 'Active:'  
Active: active (running) since Tue 2025-09-16 01:01:18 EDT; 31min ago
```

### 4.2.10 Bash & Scripting

38. Write a one-liner that loops over \*.txt in src/ and prints: : (Let's print number of words in the files)

```
annanyajain@ubuntu:~/lx-test$ for f in src/*.txt; do echo "$f: $(wc -w <  
→ "$f")";done  
src/file1.txt: 4  
src/file2.txt: 4
```

39. Write a command that exports CSV rows where dept == "CS" to cs.txt (exclude header).

So here, let me first see the structure of a csv file in src directory:

```
annanyajain@ubuntu:~/lx-test$ cat people.csv  
id,name,dept  
1,Ada,EE  
2,Linus,CS  
3,Grace,EE  
4,Dennis,CS
```

Now I know that third column is the dept:

```
annanyajain@ubuntu:~/lx-test$ awk -F, 'NR>1 && $3=="CS" {print}' people.csv >  
→ cs.txt
```

The output is redirected into the cs.txt file. We can now verify the results:

```
annanyajain@ubuntu:~/lx-test$ cat cs.txt  
2,Linus,CS  
4,Dennis,CS
```

40. Create a variable X with value 42, print it, then remove it from the environment.

```
annanyajain@ubuntu:~/lx-test$ export X=42; echo $X; unset X  
42
```

The variable `X` was created with value 42, and printed. The bash command (`unset X`) was used to remove it from the environment. Now, let's verify whether the variable is still there:

```
annanyajain@ubuntu:~/lx-test$ echo $X;
```

Nothing got printed. Hence, it is removed from the environment.

# Chapter 5

## Project Proposal

– Luo Xu, Annanya Jain, Gavin Lam

Our project is an AI Health Voice Assistant named AVA, short for Artificial Voice Assistant. Users will be able to log in and create an account and chat with AVA. This project aims to support users in tracking emotions, managing reminders, and accessing mental wellness resources. Unlike traditional chatbots, our assistant goes beyond simple question–answer interaction by integrating:

- Agentic AI workflows, enabling the assistant to interpret user intent, plan actions, and decide between generating responses, retrieving wellness exercises, or scheduling reminders.
- DevSecOps practices to make it scalable, and secure from development to deployment. The end goal is a functional prototype that not only demonstrates AI-driven health support but also serves as a practical use of modern DevOps pipelines and monitoring

This project will be an enhancement of our previously created project for CS555 course. Since then we have grown a lot in terms of knowledge and skill sets and we would like to improve it to have agentic powers to better aid users. Here is the link for the repository for the project: <https://github.com/cascadingluo/SSW590-team-7-project>

Tool	Usage
Flask	Web framework.
MongoDB	Database for user information and chat history storage.
Gemini	AI API used for the chatbot.
GitHub	Source Control and Collaboration
CI/CD	Github Actions for automated testing
Jira	Issue tracking and Agile project management tool.
AWS	Pushing local Docker with AWS and deploying with App Runner.
TikZ	Vizualizing Architecture of our assistant.

# Chapter 6

## AWS

– Gavin Lam, Luo Xu, Annanya Jain

To get our website deployed, we followed the steps in the DockerLocalAndAWS.pdf document with some changes to the code as parts in the document did not work on our machines.

First we created an AWS root account at [aws.amazon.com](https://aws.amazon.com). MFA was enabled on the root account and our region was set to us-east-2. A cost budget was set up as well to monitor costs. The SSO start URL was also recorded to be used later

We then enabled the IAM Identity center and created a new user with an alternate email. At first this user wasn't able to access any dashboards as it had no roles or permissions. This user was given the permission set AdministratorAccess by the root account and after a relog was able to access a dashboard. From this dashboard the AWS Access Key ID, AWS Secret Access Key, and AWS session ID were written down.

After the AWS accounts were correctly set up, we ran the commands in the document on the terminal.

```
aws configure sso
```

This command prompted us to give answers

```
#SSO session name:
#SSO start URL:
#SSO region:
#Account:
#Role:
#Default region:
#Output:
```

After entering all the details the login was saved.

The command ran next was

```
aws configure
```

This command prompted us to give answers



```
#AWS Access key ID:
#AWS Secret Access key:
#AWS Session ID:
#Default region:
#Output:
```

After entering all the details we successfully logged in and ran another command to confirm.

```
aws sts get-caller-identity
```

The next steps were to create an ECR repository as stated in the document. We first set some environment variables through these commands.

```
export AWS_REGION=us-east-2
export ECR_REPO=myapp
export IMAGE_TAG=v1
export CONTAINER_PORT=3000
export AWS_ACCOUNT_ID="$(aws sts get-caller-identity --query Account --output
↪ text --profile (my account here))"
```

After setting the variables we created a new ECR repository through these commands.

```
aws ecr create-repository \
  --repository-name "$ECR_REPO" \
  --image-scanning-configuration scanOnPush=true \
  --region "AWS_REGION" --profile (my account here)
```

AWS CLI was then used to obtain a short-lived registry token and logged in Docker.

```
aws ecr get-login-password --region "$AWS_REGION" --profile default \
| docker login --username AWS --password-stdin \
  "$AWS_ACCOUNT_ID.dkr.ecr.$AWS_REGION.amazonaws.com"
```

Our local image was then built, tagged, and pushed. These commands were run from the directory containing the Dockerfile.

```
docker build --platform linux/amd64 -t "$ECR_REPO:$IMAGE_TAG" .

docker tag "$ECR_REPO:$IMAGE_TAG" \
"$AWS_ACCOUNT_ID.dkr.ecr.$AWS_REGION.amazonaws.com/$ECR_REPO:$IMAGE_TAG"

docker push
↪ "$AWS_ACCOUNT_ID.dkr.ecr.$AWS_REGION.amazonaws.com/$ECR_REPO:$IMAGE_TAG"
```

Verifying the image in ECR was next

```
aws ecr describe-images \
  --repository-name "$ECR_REPO" \
  --region "$AWS_REGION" --profile default \
  --query 'imageDetails[].imageTags'
```

The next step of quick deploying with App Runner was where we ran into a roadblock. The code shown when inputted resulted in the error below.

```
export APP_NAME=my-apprunner-app

aws apprunner create-service \
  --service-name "$APP_NAME" \
  --region "$AWS_REGION" --profile default \
  --source-configuration "{
    \"ImageRepository\": {
      \"ImageIdentifier\": \"$AWS_ACCOUNT_ID.dkr.ecr.$AWS_REGION.amazonaws.com/$ECR_REPO:$IMAGE_TAG\",
      \"ImageRepositoryType\": \"ECR\",
      \"ImageConfiguration\": {\"Port\": \"$CONTAINER_PORT\"}
    },
    \"AutoDeploymentsEnabled\": true
  }" \
  --instance-configuration "{\"Cpu\": \"1 vCPU\", \"Memory\": \"2 GB\"}"
```

An error occurred (InvalidRequestException) when calling the CreateService operation: Authentication configuration is invalid.

After looking at online resources we assumed that the permission set AdministratorAccess was not working as we hoped and didn't give us the permissions we needed to use CreateService. In order to fix this we created a trust policy named AppRunnerECRAccessRole in IAM roles. The json used is:

```
{ "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "iam:PassRole",
      "Resource": "arn:aws:iam::(our account id
        ↪ here):role/AppRunnerECRAccessRole"
    },
    {
      "Effect": "Allow",
      "Action": [ "apprunner:CreateService", "apprunner:UpdateService",
        "apprunner>DeleteService" ],
      "Resource": "*" }
  ]
}
```

After relogging the SSO session and rerunning early commands to confirm we are logged in properly everywhere we retried the original deploying command again. We did some more research and found out more commands we could try. First we had to update the trust policy to replace the second bracket of code.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": { "Service": "build.apprunner.amazonaws.com" },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

The final inline policy we used is

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "iam:PassRole",
      "Resource": "arn:aws:iam::(our account id
        ↪ here):role/AppRunnerECRAccessRole"
    },
    {
      "Effect": "Allow",
      "Principal": { "Service": "build.apprunner.amazonaws.com" },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

We then had to attach permissions to the role. The AWS managed policy `AWSAppRunnerServicePolicyForECRAccess` was added. After confirming the role we retried the command again. The error persisted. We were stuck at this point so we prompted chatGPT multiple times before it recommended us to correct the create-service command with safe quoting, i.e. use single quotes around the JSON blocks so the shell does not break the quoting. The new create-service command is

```
aws apprunner create-service \
  --service-name "$APP_NAME" \
  --region "$AWS_REGION" --profile default \
  --source-configuration '{
    "ImageRepository": {
      "ImageIdentifier": "'"$AWS_ACCOUNT_ID"'.dkr.ecr.'"$AWS_REGION"'.amazonaws.'
        ↪ com/'"$ECR_REPO"':'"$IMAGE_TAG"'",
      "ImageRepositoryType": "ECR",
    }
  }
```

```

    "ImageConfiguration": { "Port": ""$CONTAINER_PORT"" }
  },
  "AuthenticationConfiguration": {
    "AccessRoleArn":
      ↪ "arn:aws:iam::"$AWS_ACCOUNT_ID":role/AppRunnerECRAccessRole"
  },
  "AutoDeploymentsEnabled": true
}' \
--instance-configuration '{
  "Cpu": "1 vCPU",
  "Memory": "2 GB"
}'

```

This miraculously went through, and we thought we had successfully created the service. We continued following the steps in the document creating more environment variables.

```

AWS_REGION=us-east-2
PROFILE=(Our Profile here)
SERVICE_ARN=(Our Service ARN here)
}

```

We then tried to run the service status changes to running command but this resulted in another error:

```

while true; do
STATUS=$(aws apprunner describe-service \
--service-arn "$SERVICE_ARN" \
--region "$AWS_REGION" --profile "$PROFILE" \
--query 'Service.Status' --output text)
echo "Service status: $STATUS"
case "$STATUS" in RUNNING|CREATE_FAILED|OPERATION_FAILED) break ;; esac
sleep 4
done

```

An error occurred (InvalidRequestException) when calling the DescribeService  
 ↪ operation: Authentication configuration is invalid.

We then read the document again and saw that the next line was to grab the service URL and just hoped that the service was running properly.

```
aws apprunner list-services \  
--region "$AWS_REGION" --profile default \  
--query "ServiceSummaryList[?ServiceName=='$APP_NAME'].ServiceUrl" --output text
```

Unfortunately, it did not create successfully and we ran into an error website. We were stuck at this point and did not know how to proceed so we decided to leave it for a little bit and come back at a later time to retry.

After coming back at a later time we reran through the steps of, aws configure sso, aws configure (with new details this time), sanity test, resetting the environment variables (they did not change from above we just reentered them), authenticated Docker to ECR, built, tagged, and pushed the local image, verified the image in ECR, deployed with App Runner. This is where we ran into the same error, but we found a service on AWS called CloudWatch that essentially replaced the need for the "test the service status changes to running" command. CloudWatch enabled us to look at the logs of services we have tried to deploy. After checking CloudWatch we saw the previous instance of our service had failed due to the ECR repository not being recognized, and saw that our most recent creation of the service had successfully deployed. We then ran the command to get the service URL.

```
aws apprunner list-services \  
--region "$AWS_REGION" --profile default \  
--query "ServiceSummaryList[?ServiceName=='$APP_NAME'].ServiceUrl" --output text
```

Which resulted in

```
mu3fbjbbv2.us-east-2.awsapprunner.com
```

This URL works, which means our deployment with AWS was successful.

After the successful deployment, we change the code for the original website to use a class-based javascript code instead of methods. The code largely remain the same, we just moved parts of code in the html file into a javascript file and had the html file run the script instead of the code directly. We have the new ColorController class that has DOM references, event binding and the color-changing behaviors. ColorController.init() connects the app on DOMContentLoaded. Here is the code in our new javascript file:

```
class ColorController { //class  
  constructor({ blueBtnId = "blueBtn", redBtnId = "redBtn", target =  
    ↪ document.body } = {}) {
```

```

    this.target = target;
    this.blueBtn = document.getElementById(blueBtnId);
    this.redBtn = document.getElementById(redBtnId);
    this._bind();
  }

  // private method to bind event listeners, the same exact code as the ones
  // originally in html
  _bind() {
    if (this.blueBtn) {
      this.blueBtn.addEventListener("click", () => this.setColor("blue"));
    }
    if (this.redBtn) {
      this.redBtn.addEventListener("click", () => this.setColor("red"));
    }
  }

  setColor(color) {
    this.target.style.backgroundColor = color;
  }

  static init(opts) {
    return new ColorController(opts);
  }
}

document.addEventListener("DOMContentLoaded", () => {
  ColorController.init();
});

```

Here is the UML diagram for the new code (this was created with the help of GPT-5):

```

+-----+
|  ColorController  |
+-----+
| - blueBtn: Button |
| - redBtn:  Button |
| - target: Element |
+-----+
| + constructor(...) |
| - _bind(): void    |
| + setColor(c):void |
| + init(...): CC    |
+-----+

```

Notes:

- constructor wires DOM elements and calls \_bind()

- `_bind` attaches event listeners to blue/red buttons
- `setColor` updates the background color of the target
- `init` is a convenience factory method

# Chapter 7

## LaTeX Docker

– Annanya Jain, Gavin Lam, Luo Xu

In this chapter, we created a Docker container to compile a simple LaTeX document using TeX Live, which is basically what Overleaf does.

### 7.1 Steps taken for creating the docker container to compile a LaTeX File

I created the following files in a folder named: texlive-app:

```
texlive-app/  
  Dockerfile  
  main
```

#### 7.1.1 main.tex

```
\begin{document}  
  
\title{Password Policy Documentation}  
\date{\today}  
\maketitle  
Let me take an example of LaTeX file with a simple table to compile.  
\section{Password Rules}  
\begin{longtable}{|l|p{9cm}|}  
\caption{Password Rules \label{Table::Passwords}}\\  
\hline  
\textbf{User} & \textbf{Password Rule / Hint} \\  
\hline  
\endhead
```



```

devuser & At least 8 characters, include uppercase,
lowercase, a number, and a symbol. Hint: First pet.
\\
\hline

admin & Must change passwords every 90 days. Hint: Favorite City.
\\
\hline

tester & Must include word banana Hint: Popular dessert item.
\\
\hline

\end{longtable}

\end{document}

```

### 7.1.2 Dockerfile

```

FROM debian:stable-slim

RUN apt-get update && apt-get install -y \
    texlive-latex-base \
    texlive-latex-recommended \
    texlive-latex-extra \
    texlive-fonts-recommended \
    texlive-fonts-extra \
    && rm -rf /var/lib/apt/lists/*

WORKDIR /data

CMD ["pdflatex", "main.tex"]

```


### 7.1.3 Building the Docker image

To build and run the container:

```


docker build -t texlive-app .
docker run --rm -v $(pwd):/data texlive-app

```



png/texlive\_build.png

After running the commands, the output file `main.pdf` gets generated in the project folder on the machine.



png/main\_pdf generated.png

# Chapter 8

## Bugzilla

– Luo Xu, Gavin Lam, Annanya Jain

### 8.1 DigitalOcean Setup

We made an account in DigitalOcean via GitHub and received 200 credits via student package. We created an droplet with these stats: 2 GB Memory / 60 GB Disk / NYC3 - Ubuntu 24.04 (LTS) x64

### 8.2 Bugzilla

After setting up the droplet, we ssh into the droplet through the local terminal with the public ip: 174.138.68.199. Then I checked the updates, and installed the essential tools and prerequisites.

```
ssh root@174.138.68.199
apt update
apt upgrade -y
apt install -y git curl wget nano build-essential
apt install -y apache2 libapache2-mod-perl2 \
    mariadb-server mariadb-client \
    libcgi-pm-perl libdbi-perl libdbd-mysql-perl \
    libtemplate-perl libdatetime-perl libdatetime-timezone-perl \
    libemail-sender-perl libemail-mime-perl libxml-twig-perl \
    libgd-perl libjson-xs-perl libauthen-sasl-perl libnet-ldap-perl \
    libsoap-lite-perl libxmlrpc-lite-perl libtest-taint-perl \
    libhtml-scrubber-perl libfile-mimeinfo-perl libcache-memcached-perl \
    perlmagick graphviz lynx python3-sphinx
```

I then configured the database.

```
systemctl start mariadb
systemctl enable mariadb
mysql_secure_installation
```

I logged into the database.

```
mysql -u root -p
```

and then inside the shell i set up a user in the SQL shell.

```
CREATE DATABASE bugzilla;
CREATE USER 'bugzillauser'@'localhost' IDENTIFIED BY '<password here>';
GRANT ALL PRIVILEGES ON bugzilla.* TO 'bugzillauser'@'localhost';
FLUSH PRIVILEGES;
EXIT;
```

I downloaded and configured bugzilla.

```
cd /var/www
git clone https://github.com/bugzilla/bugzilla.git
# Or download a tarball, e.g. wget from bugzilla.org, then extract
```

Here was when we realized we might have made an mistake of not running this in docker, so we went to stop the services based on what chat said, and resetted the host services within the droplet.

```
sudo systemctl stop apache2 || true
sudo systemctl disable apache2 || true
sudo systemctl stop mariadb || true
sudo systemctl disable mariadb || true
```

Then we installed docker and all the compose plugins incase we are missing anything.

```
sudo apt update
sudo apt install -y ca-certificates curl gnupg
sudo install -m 0755 -d /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | \
  sudo gpg --dearmor -o /etc/apt/keyrings/docker.gpg
echo \
  "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg]
  ↪ \
  https://download.docker.com/linux/ubuntu $(. /etc/os-release; echo
  ↪ $VERSION_CODENAME) stable" | \
  sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
sudo apt update
sudo apt install -y docker-ce docker-ce-cli containerd.io docker-buildx-plugin
  ↪ docker-compose-plugin
```

Then I made a project folder and env file.

```
mkdir -p ~/bugzilla-docker
cd ~/bugzilla-docker

cat > .env << 'EOF'
# ---- DB credentials (choose your own secure values) ----
MYSQL_ROOT_PASSWORD=<password>
BUGZ_DB=bugzilla
```

```
BUGZ_USER=bugzillauser
BUGZ_PASS=<password>

# ---- Internal service names / URLs ----
BUGZ_HOST=bugzilla
# If you have a domain now, set https URL. If not, set http://<PUBLIC_IP> for now
→ and change later.
BUGZ_URL=http://<PUBLIC_IP>
EOF
```

I created docker-compose.yml.

```
services:
  db:
    image: mariadb:10.6
    restart: unless-stopped
    environment:
      MYSQL_ROOT_PASSWORD: ${MYSQL_ROOT_PASSWORD}
      MYSQL_DATABASE: ${BUGZ_DB}
      MYSQL_USER: ${BUGZ_USER}
      MYSQL_PASSWORD: ${BUGZ_PASS}
    command:
      → ["--character-set-server=utf8mb4","--collation-server=utf8mb4_unicode"]
    volumes:
      - db_data:/var/lib/mysql
    networks: [bugznet]

  bugzilla:
    image: nasqueron/bugzilla:latest
    depends_on:
      db:
        condition: service_healthy
    restart: unless-stopped
    environment:
      DB_HOST: db
      DB_USER: ${BUGZ_USER}
      DB_PASSWORD: ${BUGZ_PASS}
      DB_DATABASE: ${BUGZ_DB}
      BUGZILLA_URL: ${BUGZ_URL}
    ports:
      - "80:80"
    volumes:
      - bug_data:/var/www/html/bugzilla
    networks: [bugznet]

volumes:
  db_data:
  bug_data:
```



Figure 8.1: bugzilla website on <http://174.138.68.199/>

```
networks:  
  bugznet:
```

I create the stack and checked the logs for the email and password.

```
docker compose down && docker compose up -d  
docker compose logs -f bugzilla  
docker compose exec -it db bash  
docker compose restart bugzilla  
docker compose logs -f bugzilla  
...  
bugzilla-1 | If no admin account is already defined in your database, this one  
→ will be created:  
bugzilla-1 |  
bugzilla-1 |           E-mail ..... admin@domain.tld  
bugzilla-1 |           Password ... OdMW0br43g6VW2uG8  
...
```

I then went to the bugzilla site.

logged in with the given email and password, then head to Administration → Users and re-setted the password and email to real ones.  
and now everything is running on <http://174.138.68.199/>. :)



# Chapter 9

## Overleaf

– Annanya Jain, Luo Xu, Gavin Lam

### 9.1 Overview

This document explains the complete process of deploying **Overleaf (Community Edition)** on a DigitalOcean droplet that already hosted **Bugzilla** via Docker. It includes the initial access issues, configuration steps, and debugging process that led to a fully functional Overleaf instance running alongside Bugzilla on the same host.

### 9.2 Initial Access Issue

At the beginning, I could not connect to the droplet from my local machine using its public IP address. Both `ssh root@174.138.68.199` and web requests to `http://174.138.68.199` were failing with timeout errors.

Since I didn't originally create the droplet, I did not have SSH access permissions. To resolve this, my teammate who had created the droplet and had root access, added my GitHub SSH key to the droplet's authorized keys through the DigitalOcean dashboard:

1. Logged into **DigitalOcean**.
2. Opened the droplet's page.
3. Navigated to **Access → Add SSH Keys**.
4. Added my GitHub SSH key (fetched automatically from my GitHub account).

After this, I was able to connect successfully:

```
ssh root@174.138.68.199
```

Once connected, I could verify that Bugzilla was already running:

```
docker ps
```

Bugzilla was occupying port 80, so I decided to host Overleaf on a different port.

## 9.3 Overleaf Setup Process

### 9.3.1 Step 1: Create Project Directory

```
mkdir ~/overleaf-docker
cd ~/overleaf-docker
```

### 9.3.2 Step 2: Create Environment File

```
nano .env
```

Contents:

```
OVERLEAF_PORT=8090
OVERLEAF_DOMAIN=http://174.138.68.199:8090
OVERLEAF_ADMIN_EMAIL=(email here)
OVERLEAF_ADMIN_PASSWORD=(password here)
MONGO_INITDB_ROOT_USERNAME=(username here)
MONGO_INITDB_ROOT_PASSWORD=(password here)
```

### 9.3.3 Step 3: Write the Docker Compose File

```
1 version: "3.8"
2
3 services:
4   mongo:
5     image: mongo:6.0
6     restart: unless-stopped
7     command: ["--replSet", "rs0", "--bind_ip_all"]
8     volumes:
9       - mongo_data:/data/db
10
11   redis:
12     image: redis:7
13     restart: unless-stopped
14
15   overleaf:
16     image: sharelatex/sharelatex:latest
17     restart: unless-stopped
18     depends_on:
19       - mongo
20       - redis
21     ports:
22       - "8090:80"
23     environment:
24       OVERLEAF_MONGO_URL: mongodb://mongo:27017/sharelatex?replicaSet=rs0
```

```
25     OVERLEAF_REDIS_HOST: redis
26     OVERLEAF_SITE_URL: http://174.138.68.199:8090
27     OVERLEAF_ADMIN_EMAIL: admin@example.com
28     OVERLEAF_ADMIN_PASSWORD: password123
29     OVERLEAF_APP_NAME: Overleaf
30     OVERLEAF_ALLOW_PUBLIC_REGISTRATION: "true"
31
32 volumes:
33     mongo_data:
```

### 9.3.4 Step 4: Launch the Containers

```
docker compose up -d
docker ps
```

At this point, MongoDB, Redis, and Overleaf containers appeared, but Overleaf was continuously restarting.

## 9.4 Debugging the Overleaf Restart Loop

### 9.4.1 Step 1: Check Logs

```
docker logs overleaf-docker-overleaf-1 --tail 40
```

### 9.4.2 Step 2: Error Observed


```
The MongoDB server has featureCompatibilityVersion=5.0,
but Overleaf requires at least version 6.0.
Aborting.
```

This indicated that MongoDB was version 6.0, but its internal `featureCompatibilityVersion` (FCV) was still set to 5.0.

## 9.5 Fixing MongoDB Configuration

### 9.5.1 Step 1: Enter Mongo Shell

```
docker exec -it overleaf-docker-mongo-1 mongosh
```



png/restarting\_problem\_overleaf.png

### 9.5.2 Step 2: Initialize Replica Set

```
rs.initiate()
```

If you see `MongoServerError[AlreadyInitialized]`, it means it's already set up — that's fine.

### 9.5.3 Step 3: Update Feature Compatibility Version

```
use admin
db.adminCommand({ setFeatureCompatibilityVersion: "6.0" })
```

Expected output:

```
{ "ok" : 1 }
```

### 9.5.4 Step 4: Restart Overleaf

```
exit
docker restart overleaf-docker-overleaf-1
```

## 9.6 Making Overleaf Publicly Accessible

Since Bugzilla was already bound to port 80, Overleaf was assigned to port 8090. I confirmed the port was open:

```
ss -tuln | grep 8090
```

Then enabled it in the firewall:

```
sudo ufw allow 8090/tcp
sudo ufw reload
```

## 9.7 Access and Verification

Once restarted, Overleaf was reachable at:

```
http://174.138.68.199:8090
```

I verified it via:

```
curl -I http://174.138.68.199:8090
```

Result:

```
HTTP/1.1 200 OK
```

## 9.8 Conclusion

The deployment succeeded after:

1. Gaining SSH access by adding my GitHub SSH key to the droplet.
2. Running Overleaf on port 8090 (to avoid conflict with Bugzilla on port 80).
3. Initializing the MongoDB replica set.
4. Updating the MongoDB feature compatibility version to 6.0.

After these steps, Overleaf was fully functional and accessible publicly via:

<http://174.138.68.199:8090>

This was what was seen:  
this is a test


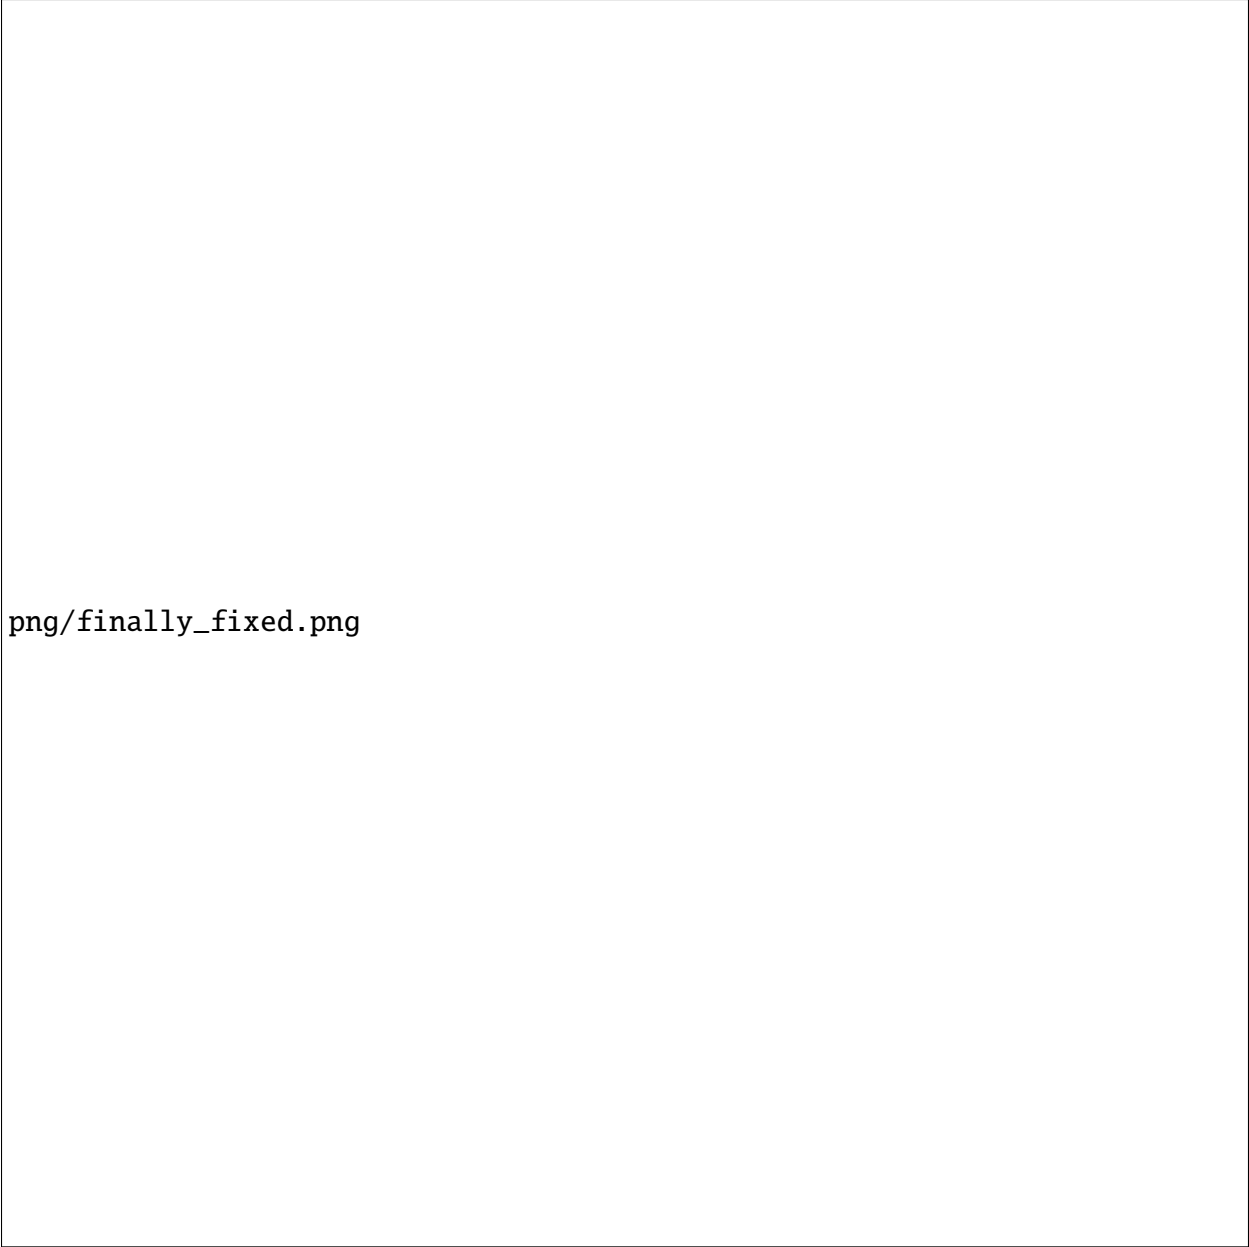


Figure 9.1: overleaf website on <http://174.138.68.199:8090>



png/finally\_fixed.png



# Appendix A

## Appendix

– *Gavin Lam, Spurthi Setty, Annanya Jain, Luo Xu*



# Bibliography

# Index

appendix, [41](#)

AWS, [16](#)

Bugzilla, [28](#)

Chapter

Appendix, [41](#)

AWS, [16](#)

Bugzilla, [28](#)

Hosts, [2](#)

Kanban Setup, [3](#)

LaTeX Docker, [24](#)

Linux Commands, [7](#)

Overleaf, [33](#)

Passwords, [1](#)

Project Proposal, [15](#)

Hosts, [2](#)

Kanban Setup, [3](#)

LaTeX Docker, [24](#)

Linux Commands, [7](#)

Overleaf, [33](#)

Passwords, [1](#)

Project Proposal, [15](#)